

Modelling Count Variables with R

Zero Truncated Poisson Distribution

Zero-Truncated Poisson Regression

- ▶ Zero-truncated Modelling is used to model count data for which the value zero cannot occur.

Examples of Zero-Truncated Model

Example 1.

- ▶ A study of length of hospital stay, in days, as a function of age, kind of health insurance and whether or not the patient died while in the hospital.
- ▶ Length of hospital stay is recorded as a minimum of at least one day.

Examples of Zero-Truncated Model

Example 2.

- ▶ A study of the number of journal articles published by tenured faculty as a function of discipline (fine arts, science, social science, humanities, medical, etc).
- ▶ To get tenure faculty must publish, therefore, there are no tenured faculty with zero publications.

Examples of Zero-Truncated Model

Example 3.

- ▶ A study by the county traffic court on the number of tickets received by teenagers as predicted by school performance, amount of driver training and gender.
- ▶ Only individuals who have received at least one citation are in the traffic court files.

Examples of Zero-Truncated Model

Example 4.

- ▶ Consider for example the random variable of the number of items in a shopper's basket at a supermarket checkout line.
- ▶ Presumably a shopper does not stand in line with nothing to buy (i.e. the minimum purchase is 1 item), so this phenomenon may follow a ZTP distribution

Zero-Truncated Poisson regression


Data Set : **hospitalstay**

- ▶ We have a hypothetical data file, **hospitalstay** with 1,493 observations.
- ▶ The length of hospital stay variable is **stay**.
- ▶ The variable **age** gives the age group from 1 to 9 which will be treated as interval in this example.
- ▶ The variables **hmo** and **died** are binary indicator variables for HMO insured patients and patients who died while in the hospital, respectively.

Zero-Truncated Poisson regression

Data Set : hospitalstay

##	stay	age	hmo	died
##	Min. : 1.00	Min. :1.00	0:1254	0:981
##	1st Qu.: 4.00	1st Qu.:4.00	1: 239	1:512
##	Median : 8.00	Median :5.00		
##	Mean : 9.73	Mean :5.23		
##	3rd Qu.:13.00	3rd Qu.:6.00		
##	Max. :74.00	Max. :9.00		



Zero-Truncated Poisson regression

Data Set : hospitalstay

- ▶ For the lowest ages, a smaller proportion of people in HMOs died, but for higher ages, there does not seem to be a huge difference, with a slightly higher proportion in HMOs dying if anything.
- ▶ Overall, as age group increases, the proportion of those dying increases, as expected.

Zero-Truncated Poisson regression

- ▶ To fit the zero-truncated Poisson model, we use the `vglm` function in the VGAM package.
- ▶ This function fits a very flexible class of models called **vector generalized linear models** to a wide range of assumed distributions.
- ▶ In our case, we believe the data are Poisson, but without zeros.
- ▶ Thus the values are strictly positive Poisson, for which we use the positive Poisson family via the `pospoisson` function passed to `vglm`.

Zero-Truncated Poisson regression

Fitting the Model with R

We will use the *hospitalstay* data.

```
m1 <- vglm(stay ~ age + hmo + died,  
            family = pospoisson(),  
            data = hospitalstay)  
summary(m1)
```

Zero-Truncated Poisson regression

Fitting the Model with R

Model Summary

Coefficients:


##	Estimate	Std. Error	z	value
## (Intercept)	2.436	0.027	89.1	
## age	-0.014	0.005	-2.9	
## hmo1	-0.136	0.024	-5.7	
## died1	-0.204	0.018	-11.1	

Zero-Truncated Poisson regression

- ▶ The value of the coefficient for age, -0.0144 suggests that the log count of stay decreases by 0.0144 for each year increase in age.
- ▶ The coefficient for hmo, -0.1359 indicates that the log count of stay for HMO patient is 0.1359 less than for non-HMO patients.
- ▶ The log count of stay for patients who died while in the hospital was 0.2038 less than those patients who did not die.
- ▶ Finally, the value of the constant 2.4358 is the log count of the stay when all of the predictors equal zero.

Zero-Truncated Poisson regression

- ▶ Can compute CIs using **boot** package
- ▶ Age does not have a significant effect, but hmo and died both do.



Zero-truncated negative binomial regression

- ▶ Zero-truncated negative binomial regression is used to model count data for which the value zero cannot occur and for which over dispersion exists.

Zero-truncated negative binomial regression

- ▶ To fit the zero-truncated negative binomial model, we use the `vglm` function in the VGAM package.
- ▶ This function fits a very flexible class of models called vector generalized linear models to a wide range of assumed distributions.
- ▶ In our case, we believe the data come from the negative binomial distribution, but without zeros.
- ▶ Thus the values are strictly positive poisson, for which we use the positive negative binomial family via the `posnegbinomial` function passed to `vglm`.

Zero-truncated negative binomial regression

Fitting the Model with R

We will use the *hospitalstay* data again.

```
m1 <- vglm(stay ~ age + hmo + died,  
  family = posnegbinomial(),  
  data = hospitalstay)
```

Zero Truncated Negative Binomial Regression

```
summary(m1)
##
## Call:
## vglm(formula = stay ~ age + hmo + died,
##       family = posnegbinomial(),
##       data = hospitalstay)
##
## Pearson Residuals:
##           Min      1Q  Median     3Q      Max
## log(munb)  -1.4 -0.70   -0.23  0.45  9.8
## log(size) -14.1 -0.27    0.45  0.76  1.0
```

Zero Truncated Negative Binomial Regression

Coefficients:

##	Estimate	Std. Error	z	value
## (Intercept):1	2.408	0.072	33.6	
## (Intercept):2	0.569	0.055	10.4	
## age	-0.016	0.013	-1.2	
## hmo1	-0.147	0.059	-2.5	
## died1	-0.218	0.046	-4.7	

Zero Truncated Negative Binomial Regression

- ▶
- ▶ The first intercept is what we know as the typical intercept.
- ▶ The second is the **over dispersion parameter**, α .
- ▶ The number of linear predictors is two, one for the expected mean λ and one for the over dispersion.
- ▶ Next the dispersion parameter is printed, assumed to be one after accounting for overdispersion.

Zero Truncated Negative Binomial Regression

- ▶ The value of the coefficient for age, -0.0157 suggests that the log count of stay decreases by 0.0157 for each year increase in age.
- ▶ The coefficient for hmo, -0.1471 indicates that the log count of stay for HMO patient is 0.1471 less than for non-HMO patients.
- ▶ The log count of stay for patients who died while in the hospital was 0.2178 less than those patients who did not die.

Zero Truncated Negative Binomial Regression

- ▶ The value of the constant 2.4083 is the log count of the stay when all of the predictors equal zero.
- ▶ The value of the second intercept, the over dispersion parameter, α is 0.5686.
- ▶ To test whether we need to estimate over dispersion, we could fit a zero-truncated Poisson model and compare the two. (Not Covered).